

sa-20231003-1696291200-jupyterlite

June 20, 2024

Final Project: House Sales in King County, USA

Table of Contents

Instructions

About the Dataset

Module 1: Importing Data

Module 2: Data Wrangling

Module 3: Exploratory Data Analysis

Module 4: Model Development

Module 5: Model Evaluation and Refinement

Estimated Time Needed: 75 min

1 Instructions

In this assignment, you are a Data Analyst working at a Real Estate Investment Trust. The Trust would like to start investing in Residential real estate. You are tasked with determining the market price of a house given a set of features. You will analyze and predict housing prices using attributes or features such as square footage, number of bedrooms, number of floors, and so on. This is a template notebook; your job is to complete the ten questions. Some hints to the questions are given.

As you are completing this notebook, take and save the **screenshots** of the final outputs of your solutions (e.g., final charts, tables, calculation results etc.). They will need to be shared in the following Peer Review section of the Final Project module.

2 About the Dataset

This dataset contains house sale prices for King County, which includes Seattle. It includes homes sold between May 2014 and May 2015. It was taken from [here](#). It was also slightly modified for the purposes of this course.

Variable	Description
id	A notation for a house
date	Date house was sold
price	Price is prediction target
bedrooms	Number of bedrooms

Variable	Description
bathrooms	Number of bathrooms
sqft_living	Square footage of the home
sqft_lot	Square footage of the lot
floors	Total floors (levels) in house
waterfront	House which has a view to a waterfront
view	Has been viewed
condition	How good the condition is overall
grade	overall grade given to the housing unit, based on King County grading system
sqft_above	Square footage of house apart from basement
sqft_basement	Square footage of the basement
yr_built	Built Year
yr_renovated	Year when house was renovated
zipcode	Zip code
lat	Latitude coordinate
long	Longitude coordinate
sqft_living15	Living room area in 2015(implies– some renovations) This might or might not have affected the lotsize area
sqft_lot15	LotSize area in 2015(implies– some renovations)

2.1 Import the required libraries

```
[14]: # All Libraries required for this lab are listed below. The libraries
      ↪pre-installed on Skills Network Labs are commented.
      # !mamba install -qy pandas==1.3.4 numpy==1.21.4 seaborn==0.9.0 matplotlib==3.5.
      ↪0 scikit-learn==0.20.1
      # Note: If your environment doesn't support "!mamba install", use "!pip install"
```

```
[15]: # Surpress warnings:
def warn(*args, **kwargs):
    pass
import warnings
warnings.warn = warn
```

```
[16]: #!pip install -U scikit-learn
```

```
[47]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler,PolynomialFeatures
from sklearn.linear_model import LinearRegression
%matplotlib inline
```

3 Module 1: Importing Data Sets

Download the dataset by running the cell below.

```
[18]: import piplite
      await piplite.install('seaborn')
```

```
[19]: from pyodide.http import pyfetch

      async def download(url, filename):
          response = await pyfetch(url)
          if response.status == 200:
              with open(filename, "wb") as f:
                  f.write(await response.bytes())
```

```
[20]: filepath='https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
      ↪IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/FinalModule_Coursera/
      ↪data/kc_house_data_NaN.csv'
```

```
[21]: await download(filepath, "housing.csv")
      file_name="housing.csv"
```

Load the csv:

```
[22]: df = pd.read_csv(file_name)
```

Note: This version of the lab is working on JupyterLite, which requires the dataset to be downloaded to the interface. While working on the downloaded version of this notebook on their local machines (Jupyter Anaconda), the learners can simply **skip the steps above**, and simply use the URL directly in the `pandas.read_csv()` function. You can uncomment and run the statements in the cell below.

```
[23]: #filepath='https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
      ↪IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/FinalModule_Coursera/
      ↪data/kc_house_data_NaN.csv'
      #df = pd.read_csv(filepath, header=None)
```

We use the method `head` to display the first 5 columns of the dataframe.

```
[24]: df.head()
```

```
[24]: Unnamed: 0      id      date      price  bedrooms  bathrooms  \
0         0  7129300520  20141013T000000  221900.0         3.0         1.00
1         1  6414100192  20141209T000000  538000.0         3.0         2.25
2         2  5631500400  20150225T000000  180000.0         2.0         1.00
3         3  2487200875  20141209T000000  604000.0         4.0         3.00
4         4  1954400510  20150218T000000  510000.0         3.0         2.00
```

```
sqft_living  sqft_lot  floors  waterfront  ...  grade  sqft_above  \
```

0	1180	5650	1.0	0 ...	7	1180
1	2570	7242	2.0	0 ...	7	2170
2	770	10000	1.0	0 ...	6	770
3	1960	5000	1.0	0 ...	7	1050
4	1680	8080	1.0	0 ...	8	1680

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long \
0	0	1955	0	98178	47.5112	-122.257
1	400	1951	1991	98125	47.7210	-122.319
2	0	1933	0	98028	47.7379	-122.233
3	910	1965	0	98136	47.5208	-122.393
4	0	1987	0	98074	47.6168	-122.045

	sqft_living15	sqft_lot15
0	1340	5650
1	1690	7639
2	2720	8062
3	1360	5000
4	1800	7503

[5 rows x 22 columns]

3.0.1 Question 1

Display the data types of each column using the function `dtypes`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[25]: # Display data types
print(df.dtypes)
```

```
Unnamed: 0      int64
id              int64
date            object
price           float64
bedrooms        float64
bathrooms        float64
sqft_living      int64
sqft_lot         int64
floors           float64
waterfront       int64
view             int64
condition        int64
grade            int64
sqft_above       int64
sqft_basement    int64
yr_built         int64
yr_renovated     int64
zipcode          int64
```

```

lat                float64
long               float64
sqft_living15      int64
sqft_lot15         int64
dtype: object

```

We use the method describe to obtain a statistical summary of the dataframe.

```
[26]: df.describe()
```

```

[26]:      Unnamed: 0      id      price      bedrooms      bathrooms \
count  21613.00000  2.161300e+04  2.161300e+04  21600.000000  21603.000000
mean    10806.00000  4.580302e+09  5.400881e+05      3.372870      2.115736
std      6239.28002  2.876566e+09  3.671272e+05      0.926657      0.768996
min         0.00000  1.000102e+06  7.500000e+04      1.000000      0.500000
25%      5403.00000  2.123049e+09  3.219500e+05      3.000000      1.750000
50%     10806.00000  3.904930e+09  4.500000e+05      3.000000      2.250000
75%     16209.00000  7.308900e+09  6.450000e+05      4.000000      2.500000
max     21612.00000  9.900000e+09  7.700000e+06     33.000000      8.000000

      sqft_living      sqft_lot      floors      waterfront      view \
count  21613.000000  2.161300e+04  21613.000000  21613.000000  21613.000000
mean    2079.899736  1.510697e+04      1.494309      0.007542      0.234303
std      918.440897  4.142051e+04      0.539989      0.086517      0.766318
min      290.000000  5.200000e+02      1.000000      0.000000      0.000000
25%     1427.000000  5.040000e+03      1.000000      0.000000      0.000000
50%     1910.000000  7.618000e+03      1.500000      0.000000      0.000000
75%     2550.000000  1.068800e+04      2.000000      0.000000      0.000000
max     13540.00000  1.651359e+06      3.500000      1.000000      4.000000

      ...      grade      sqft_above      sqft_basement      yr_built \
count  ...  21613.000000  21613.000000  21613.000000  21613.000000
mean    ...      7.656873  1788.390691      291.509045  1971.005136
std      ...      1.175459   828.090978   442.575043    29.373411
min      ...      1.000000   290.000000      0.000000  1900.000000
25%      ...      7.000000  1190.000000      0.000000  1951.000000
50%      ...      7.000000  1560.000000      0.000000  1975.000000
75%      ...      8.000000  2210.000000     560.000000  1997.000000
max      ...     13.000000  9410.000000   4820.000000  2015.000000

      yr_renovated      zipcode      lat      long      sqft_living15 \
count  21613.000000  21613.000000  21613.000000  21613.000000  21613.000000
mean      84.402258  98077.939805     47.560053   -122.213896  1986.552492
std     401.679240    53.505026     0.138564     0.140828   685.391304
min         0.000000  98001.000000     47.155900   -122.519000   399.000000
25%         0.000000  98033.000000     47.471000   -122.328000  1490.000000
50%         0.000000  98065.000000     47.571800   -122.230000  1840.000000
75%         0.000000  98118.000000     47.678000   -122.125000  2360.000000

```

max	2015.000000	98199.000000	47.777600	-121.315000	6210.000000
-----	-------------	--------------	-----------	-------------	-------------

	sqft_lot15
count	21613.000000
mean	12768.455652
std	27304.179631
min	651.000000
25%	5100.000000
50%	7620.000000
75%	10083.000000
max	871200.000000

[8 rows x 21 columns]

4 Module 2: Data Wrangling

4.0.1 Question 2

Drop the columns “id” and “Unnamed: 0” from axis 1 using the method `drop()`, then use the method `describe()` to obtain a statistical summary of the data. Make sure the `inplace` parameter is set to `True`. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.

```
[27]: #Enter Your Code, Execute and take the Screenshot
# Drop columns 'id' and 'Unnamed: 0'
df.drop(columns=['id', 'Unnamed: 0'], inplace=True)

# Get summary statistics
print(df.describe())
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot \
count	2.161300e+04	21600.000000	21603.000000	21613.000000	2.161300e+04
mean	5.400881e+05	3.372870	2.115736	2079.899736	1.510697e+04
std	3.671272e+05	0.926657	0.768996	918.440897	4.142051e+04
min	7.500000e+04	1.000000	0.500000	290.000000	5.200000e+02
25%	3.219500e+05	3.000000	1.750000	1427.000000	5.040000e+03
50%	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	6.450000e+05	4.000000	2.500000	2550.000000	1.068800e+04
max	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

	floors	waterfront	view	condition	grade \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1.494309	0.007542	0.234303	3.409430	7.656873
std	0.539989	0.086517	0.766318	0.650743	1.175459
min	1.000000	0.000000	0.000000	1.000000	1.000000
25%	1.000000	0.000000	0.000000	3.000000	7.000000
50%	1.500000	0.000000	0.000000	3.000000	7.000000

75%	2.000000	0.000000	0.000000	4.000000	8.000000
max	3.500000	1.000000	4.000000	5.000000	13.000000

	sqft_above	sqft_basement	yr_built	yr_renovated	zipcode \
count	21613.000000	21613.000000	21613.000000	21613.000000	21613.000000
mean	1788.390691	291.509045	1971.005136	84.402258	98077.939805
std	828.090978	442.575043	29.373411	401.679240	53.505026
min	290.000000	0.000000	1900.000000	0.000000	98001.000000
25%	1190.000000	0.000000	1951.000000	0.000000	98033.000000
50%	1560.000000	0.000000	1975.000000	0.000000	98065.000000
75%	2210.000000	560.000000	1997.000000	0.000000	98118.000000
max	9410.000000	4820.000000	2015.000000	2015.000000	98199.000000

	lat	long	sqft_living15	sqft_lot15
count	21613.000000	21613.000000	21613.000000	21613.000000
mean	47.560053	-122.213896	1986.552492	12768.455652
std	0.138564	0.140828	685.391304	27304.179631
min	47.155900	-122.519000	399.000000	651.000000
25%	47.471000	-122.328000	1490.000000	5100.000000
50%	47.571800	-122.230000	1840.000000	7620.000000
75%	47.678000	-122.125000	2360.000000	10083.000000
max	47.777600	-121.315000	6210.000000	871200.000000

We can see we have missing values for the columns bedrooms and bathrooms

```
[28]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
      ↪sum())
      print("number of NaN values for the column bathrooms :", df['bathrooms'].
      ↪isnull().sum())
```

```
number of NaN values for the column bedrooms : 13
number of NaN values for the column bathrooms : 10
```

We can replace the missing values of the column 'bedrooms' with the mean of the column 'bedrooms' using the method replace(). Don't forget to set the inplace parameter to True

```
[29]: mean=df['bedrooms'].mean()
      df['bedrooms'].replace(np.nan,mean, inplace=True)
```

We also replace the missing values of the column 'bathrooms' with the mean of the column 'bathrooms' using the method replace(). Don't forget to set the inplace parameter to True

```
[30]: mean=df['bathrooms'].mean()
      df['bathrooms'].replace(np.nan,mean, inplace=True)
```

```
[31]: print("number of NaN values for the column bedrooms :", df['bedrooms'].isnull().
      ↪sum())
      print("number of NaN values for the column bathrooms :", df['bathrooms'].
      ↪isnull().sum())
```

number of NaN values for the column bedrooms : 0
number of NaN values for the column bathrooms : 0

5 Module 3: Exploratory Data Analysis

5.0.1 Question 3

Use the method `value_counts` to count the number of houses with unique floor values, use the method `.to_frame()` to convert it to a data frame. Take a screenshot of your code and output. You will need to submit the screenshot for the final project.


```
[32]: #Enter Your Code, Execute and take the Screenshot  
# Count unique values in the 'floors' column and convert to DataFrame  
floor_counts = df['floors'].value_counts().to_frame()  
  
# Display result  
print(floor_counts)
```

```
      floors  
1.0    10680  
2.0     8241  
1.5     1910  
3.0      613  
2.5      161  
3.5         8
```

5.0.2 Question 4

Use the function `boxplot` in the `seaborn` library to determine whether houses with a waterfront view or without a waterfront view have more price outliers. Take a screenshot of your code and boxplot. You will need to submit the screenshot for the final project.

```
[33]: # Import seaborn library (assuming it's not already imported)  
import seaborn as sns  
  
# Create a regression plot  
sns.regplot(x="sqft_above", y="price", data=df)  
  
# Display the plot  
plt.show()
```


5.0.3 Question 5

Use the function `regplot` in the `seaborn` library to determine if the feature `sqft_above` is negatively or positively correlated with `price`. Take a screenshot of your code and scatterplot. You will need to submit the screenshot for the final project.

```
[34]: # Import seaborn library (assuming it's not already imported)
import seaborn as sns

# Create a regression plot
sns.regplot(x="sqft_above", y="price", data=df)

# Display the plot
plt.show()
```



We can use the Pandas method `corr()` to find the feature other than price that is most correlated with price.

```
[35]: df.corr()['price'].sort_values()
```

```
[35]: zipcode      -0.053203
      long         0.021626
      condition    0.036362
      yr_built     0.054012
      sqft_lot15    0.082447
      sqft_lot      0.089661
      yr_renovated  0.126434
      floors       0.256794
      waterfront   0.266369
      lat          0.307003
      bedrooms     0.308797
      sqft_basement 0.323816
      view         0.397293
      bathrooms    0.525738
      sqft_living15 0.585379
```

```
sqft_above      0.605567
grade           0.667434
sqft_living     0.702035
price           1.000000
Name: price, dtype: float64
```

6 Module 4: Model Development

We can Fit a linear regression model using the longitude feature 'long' and caculate the R^2 .

```
[36]: X = df[['long']]
      Y = df['price']
      lm = LinearRegression()
      lm.fit(X,Y)
      lm.score(X, Y)
```

```
[36]: 0.00046769430149007363
```

6.0.1 Question 6

Fit a linear regression model to predict the 'price' using the feature 'sqft_living' then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[37]: #Enter Your Code, Execute and take the Screenshot
      from sklearn.linear_model import LinearRegression

      # Define features and target
      X = df[['sqft_living']] # sqft_living as feature
      y = df['price'] # price as target

      # Create and fit the model
      model = LinearRegression()
      model.fit(X, y)

      # Calculate R^2
      r2 = model.score(X, y)
      print(f"R^2: {r2:.2f}")
```

```
R^2: 0.49
```

6.0.2 Question 7

Fit a linear regression model to predict the 'price' using the list of features:

```
[38]: features = ["floors", "waterfront","lat" ,"bedrooms" ,"sqft_basement" ,"view"␣
      ↪,"bathrooms","sqft_living15","sqft_above","grade","sqft_living"]
```

Then calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[39]: # Check for missing values
print(df.isnull().sum())
```

```
date          0
price         0
bedrooms      0
bathrooms     0
sqft_living   0
sqft_lot      0
floors        0
waterfront    0
view          0
condition     0
grade         0
sqft_above    0
sqft_basement 0
yr_built      0
yr_renovated  0
zipcode       0
lat           0
long          0
sqft_living15 0
sqft_lot15    0
dtype: int64
```

6.0.3 This will help with Question 8

Create a list of tuples, the first element in the tuple contains the name of the estimator:

‘scale’

‘polynomial’

‘model’

The second element in the tuple contains the model constructor

StandardScaler()

PolynomialFeatures(include_bias=False)

LinearRegression()

```
[40]: Input=[('scale',StandardScaler()),('polynomial',
↳PolynomialFeatures(include_bias=False)),('model',LinearRegression())]
```

6.0.4 Question 8

Use the list to create a pipeline object to predict the 'price', fit the object using the features in the list features, and calculate the R^2 . Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[41]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import StandardScaler
      from sklearn.linear_model import LinearRegression
      from sklearn.preprocessing import PolynomialFeatures

      # Assuming you've handled missing values (refer to previous explanation)

      # Define steps
      Input = [('scale', StandardScaler()), ('polynomial',
      ↪PolynomialFeatures(include_bias=False)), ('model', LinearRegression())]

      # Create pipeline
      pipe = Pipeline(Input)

      # Fit the pipeline
      pipe.fit(X, y)

      # Get  $R^2$  score
      r2 = pipe.score(X, y)
      print(f" $R^2$ : {r2:.2f}")
```

R^2 : 0.53

7 Module 5: Model Evaluation and Refinement

Import the necessary modules:

```
[42]: from sklearn.model_selection import cross_val_score
      from sklearn.model_selection import train_test_split
      print("done")
```

done

We will split the data into training and testing sets:

```
[43]: features = ["floors", "waterfront", "lat", "bedrooms", "sqft_basement", "view",
      ↪ "bathrooms", "sqft_living15", "sqft_above", "grade", "sqft_living"]
      X = df[features]
      Y = df['price']

      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.15,
      ↪ random_state=1)
```

```
print("number of test samples:", x_test.shape[0])
print("number of training samples:", x_train.shape[0])
```

number of test samples: 3242
number of training samples: 18371

7.0.1 Question 9

Create and fit a Ridge regression object using the training data, set the regularization parameter to 0.1, and calculate the R^2 using the test data. Take a screenshot of your code and the value of the R^2 . You will need to submit it for the final project.

```
[44]: from sklearn.linear_model import Ridge
```

```
[45]: #Enter Your Code, Execute and take the Screenshot
from sklearn.linear_model import Ridge
from sklearn.model_selection import train_test_split

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪ random_state=42)

# Create Ridge regression object with alpha (regularization parameter) set to 0.
    ↪ 1
ridge_reg = Ridge(alpha=0.1)

# Fit the model on the training data
ridge_reg.fit(X_train, y_train)

# Make predictions on the testing set
y_pred = ridge_reg.predict(X_test)

# Calculate R^2 score on the testing set
from sklearn.metrics import r2_score
r2 = r2_score(y_test, y_pred)
print(f"R^2 on testing set: {r2:.2f}")
```

R^2 on testing set: 0.66

7.0.2 Question 10

Perform a second order polynomial transform on both the training data and testing data. Create and fit a Ridge regression object using the training data, set the regularisation parameter to 0.1, and calculate the R^2 utilising the test data provided. Take a screenshot of your code and the R^2 . You will need to submit it for the final project.

```
[48]: #Enter Your Code, Execute and take the Screenshot
from sklearn.linear_model import Ridge
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import PolynomialFeatures
from sklearn.model_selection import train_test_split
from sklearn.metrics import r2_score

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Create a polynomial features object for second-order terms (degree=2)
poly = PolynomialFeatures(degree=2, include_bias=False)

# Apply polynomial transform to training and testing features
X_train_poly = poly.fit_transform(X_train)
X_test_poly = poly.transform(X_test)

# Create a standard scaler object
scaler = StandardScaler()

# Scale the training and testing features (after polynomial transform)
X_train_scaled = scaler.fit_transform(X_train_poly)
X_test_scaled = scaler.transform(X_test_poly)

# Create a Ridge regression object with alpha (regularization parameter) set to
    ↪0.1
ridge_reg = Ridge(alpha=0.1)

# Fit the model on the scaled training data
ridge_reg.fit(X_train_scaled, y_train)

# Make predictions on the scaled testing set
y_pred = ridge_reg.predict(X_test_scaled)

# Calculate R2 score on the testing set
r2 = r2_score(y_test, y_pred)
print(f"R2 on testing set: {r2:.2f}")
```

R² on testing set: 0.70

Once you complete your notebook you will have to share it. You can download the notebook by navigating to “File” and clicking on “Download” button.

<p></p>

This will save the (.ipynb) file on your computer. Once saved, you can upload this file in the “My Submission” tab, of the “Peer-graded Assignment” section.

About the Authors:

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: Michelle Carey, Mavis Zhou

7.1 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-12-01	2.2	Aije Egwaikhide	Coverted Data describtion from text to table
2020-10-06	2.1	Lakshmi Holla	Changed markdown instruction of Question1
2020-08-27	2.0	Malika Singla	Added lab to GitLab
2022-06-13	2.3	Svitlana Kramar	Updated Notebook sharing instructions

|

|||

##

© IBM Corporation 2020. All rights reserved.

[]: